

VISVESVARAYA TECHNOLOGICAL UNIVERSITY

“JnanaSangama”, Belgaum -590014, Karnataka.



LAB REPORT
on

Machine Learning (23CS6PCMAL)

Submitted by

ARJUN PRABHAKARAN (1BM22CS053)

in partial fulfillment for the award of the degree of

BACHELOR OF ENGINEERING
in
COMPUTER SCIENCE AND ENGINEERING



B.M.S. COLLEGE OF ENGINEERING
(Autonomous Institution under VTU)
BENGALURU-560019
Sep-2024 to Jan-2025

B.M.S. College of Engineering,
Bull Temple Road, Bangalore 560019
(Affiliated To Visvesvaraya Technological University, Belgaum)
Department of Computer Science and Engineering



CERTIFICATE

This is to certify that the Lab work entitled “Machine Learning (23CS6PCMAL)” carried out by **ARJUN PRABHAKARAN (1BM22CS053)**, who is bonafide student of **B.M.S. College of Engineering**. It is in partial fulfillment for the award of **Bachelor of Engineering in Computer Science and Engineering** of the Visvesvaraya Technological University, Belgaum. The Lab report has been approved as it satisfies the academic requirements in respect of an Machine Learning (23CS6PCMAL) work prescribed for the said degree.

Sarala D V Assistant Professor Department of CSE, BMSCE	Dr. Kavitha Sooda Professor & HOD Department of CSE, BMSCE
---	--

Index

Sl. No.	Date	Experiment Title	Page No.
1	4-3-2025	Write a python program to import and export data using Pandas library functions	4
2	4-3-2025	Demonstrate various data pre-processing techniques for a given dataset	7
3	11-3-2025	Implement Linear and Multi-Linear Regression algorithm using appropriate dataset	13
4	18-3-2025	Build Logistic Regression Model for a given dataset	23
5	25-3-2025	Use an appropriate data set for building the decision tree (ID3) and apply this knowledge to classify a new sample.	26
6	1-4-2025	Build KNN Classification model for a given dataset.	31
7	1-4-2025	Build Support vector machine model for a given dataset	43
8	15-4-2025	Implement Random forest ensemble method on a given dataset.	50
9	15-4-2025	Implement Boosting ensemble method on a given dataset.	53
10	30-4-2025	Build k-Means algorithm to cluster a set of data stored in a .CSV file.	56
11	30-4-2025	Implement Dimensionality reduction using Principal Component Analysis (PCA) method.	62

Github Link: https://github.com/wombat-42/ML_LAB.git

Program 1

Write a python program to import and export data using Pandas library functions

Screenshot

LAB1 - Data processing

import pandas as pd
import numpy as np

1) To load .csv file to dataframe
→ df = pd.read_csv("housing.csv")

2) To display information of all columns
⇒ print(df)

3) To display statistical information of all columns
⇒ df.describe()

4) To display unique count of values for "ocean proximity" column

5) To display which attributes (columns) in data set have missing values count greater than zero.
⇒ print(df.isnull().sum())

for i in df:
if df.isnull().sum() > 0:
print(i)
or
missing-values = df.isnull().sum()
print(missing-values)

here missing values
How did you handle them

ans. The columns did not have missing values, but I added them.

→ Diabetes.csv
columns: TOR, HDL
I handled them by using SimpleImputer method of sklearn.

In the TOR column, missing values were filled by using finding the mean of all the values present in the column TOR.

For HDL also we found mean value of all values in the HDL column & filled it in place of missing values.

→ Adult.csv
columns: capital-gain, hours-per-week
missing value in capital-gain was replaced by finding the mean of all values in capital-gain & for hours-per-week, median

2. which categorical columns did you identify in the data set? How did you encode them?

ans. Diabetes.csv
Categorical data: gender
the gender column was encoded by using ordinal encoder, which replaced M with 0.0 & F with 1.0

Adult.csv:
Categorical data: workingclass
The workingclass column was encoded by using OneHotEncoder, which substitutes & creates a column with all values of workingclass placed 1 in that column if the row belongs to that category

3. What is scaling and standardization? When would you use one over the other?

min max scaling is used to transform features to be on a similar scale.
$$x_{new} = (x - x_{min}) / (x_{max} - x_{min})$$

Standardization: transformation of features by subtracting from mean & dividing by standard deviation.
$$x_{new} = (x - \text{mean}) / \text{std.}$$

minmax scaling is used for uniform data. It is affected by outliers.
Standardisation is used when the data follows Gaussian distribution. It is not affected by outliers.

10/3/25

Code:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import OrdinalEncoder, OneHotEncoder
from sklearn.preprocessing import StandardScaler, MinMaxScaler
from scipy import stats

***Diabetes Dataset**
df=pd.read_csv('/content/Dataset of Diabetes .csv')
df.head()
df.shape
print(df.info())
# Summary statistics
print(df.describe())
missing_values=df.isnull().sum()
print(missing_values[missing_values > 0])
categorical_cols = df.select_dtypes(include=['object']).columns
print("Categorical columns identified:", categorical_cols)
if len(categorical_cols) > 0:
    df = pd.get_dummies(df, columns=categorical_cols, drop_first=True)
    print("\nDataFrame after one-hot encoding:")
    print(df.head())
else:
    print("\nNo categorical columns found in the dataset.")
from sklearn.preprocessing import MinMaxScaler, StandardScaler
import pandas as pd

numerical_cols = df.select_dtypes(include=['number']).columns

scaler = MinMaxScaler()
df_minmax = df.copy() # Create a copy to avoid modifying the original
df_minmax[numerical_cols] = scaler.fit_transform(df[numerical_cols])

scaler = StandardScaler()
df_standard = df.copy()
df_standard[numerical_cols] = scaler.fit_transform(df[numerical_cols])
print("\nDataFrame after Min-Max Scaling:")
print(df_minmax.head())
print("\nDataFrame after Standardization:")
print(df_standard.head())

***Adult Income Dataset**
df1=pd.read_csv('/content/adult.csv')
df1.head()
df1.shape
```

```

print(df1.info())
# Summary statistics
print(df.describe())
missing_values=df1.isnull().sum()
print(missing_values[missing_values > 0])
categorical_cols = df1.select_dtypes(include=['object']).columns
print("Categorical columns identified:", categorical_cols)
if len(categorical_cols) > 0:
    df1 = pd.get_dummies(df1, columns=categorical_cols, drop_first=True)
    print("\nDataFrame after one-hot encoding:")
    print(df1.head())
else:
    print("\nNo categorical columns found in the dataset.")
from sklearn.preprocessing import MinMaxScaler, StandardScaler
import pandas as pd

numerical_cols = df1.select_dtypes(include=['number']).columns

scaler = MinMaxScaler()
df_minmax = df1.copy() # Create a copy to avoid modifying the original
df_minmax[numerical_cols] = scaler.fit_transform(df1[numerical_cols])

scaler = StandardScaler()
df_standard = df1.copy()
df_standard[numerical_cols] = scaler.fit_transform(df1[numerical_cols])
print("\nDataFrame after Min-Max Scaling:")
print(df_minmax.head())
print("\nDataFrame after Standardization:")
print(df_standard.head())

```

PROGRAM 2 Demonstrate various data pre-processing techniques for a given dataset

Screenshot

10/3/2024
Monday

LAB 2

- Geographical features that are:
 - longitude
 - latitude
 - ocean proximity
- Plot a graph to show features correlation with housing price. Which feature correlates to the maximum

Ans: Median income correlates the most to the maximum

- List the features that could be combined to improve correlation & plot again to see if correlation has improved

Ans: Combine → total rooms & total bedrooms
population & households.

- List features that could be combined to improve correlation ~~total~~
- List features that need to be cleaned & demonstrate the process of cleaning

Ans: Total rooms has some missing values & needs some cleaning.
You can remove the ~~all~~ null values by filling the null values with median no of bedrooms.

• List any categorical features that need to be converted to numerical?
Yes, Ocean Proximity.

• Discuss importance of feature scaling.
It ensures all features are on a similar scale, allowing for fair comparison between them & leading to improved model performance by preventing features with larger range from dominating the learning process.

Code

```
import pandas as pd

import numpy as np

import seaborn as sns

import matplotlib.pyplot as plt

df=pd.read_csv('housing.csv')

df.head(2)

df.describe()

df.info()

sns.histplot(df['median_income'], kde=True, color='green')

sns.histplot(df['housing_median_age'])

from sklearn.model_selection import train_test_split

X = df.drop("median_house_value", axis=1)

y = df["median_house_value"]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

X = df.drop("median_house_value", axis=1)

y = df["median_house_value"]

df["income_cat"] = pd.cut(df["median_house_value"],

bins=[0, 100000, 200000, 300000, 400000, np.inf],

labels=[1, 2, 3, 4, 5])

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42,

stratify=df["income_cat"])
```



```

train_set = X_train.copy()

train_set["median_house_value"] = y_train

train_set.plot(kind="scatter", x="longitude", y="latitude", alpha=0.4,s=train_set["population"]/100,
label="population",figsize=(10,7), c="median_house_value", cmap=plt.get_cmap("jet"),

colorbar=True)

plt.legend()

numerical_columns = df.select_dtypes(include=['float64', 'int64'])

correlation_matrix = numerical_columns.corr()

print(correlation_matrix["median_house_value"].sort_values(ascending=False))

df.plot(kind="scatter", x="median_income", y="median_house_value", alpha=0.1)

# Combine 'median_income' and 'households'

df["income_households"] = df["median_income"] * df["households"]


numerical_columns = df.select_dtypes(include=['float64', 'int64'])

correlation_matrix = numerical_columns.corr()

print(correlation_matrix["median_house_value"].sort_values(ascending=False))

df.plot(kind="scatter", x="income_households", y="median_house_value", alpha=0.1)

plt.show()

missing_values = df.isnull().sum()

print(missing_values[missing_values > 0])

h=df

h.dropna(subset=["total_bedrooms"])

from sklearn.preprocessing import OneHotEncoder

df1=pd.read_csv('housing.csv')

hc=df1[["ocean_proximity"]]

```

```

encoder=OneHotEncoder()

hc_encoded=encoder.fit_transform(hc).toarray()

hc_1hot_df = pd.DataFrame(hc_encoded, columns=encoder.get_feature_names_out(hc.columns))

hc_1hot_df.head()

```

Feature scaling is crucial in machine learning for several reasons, particularly when using algorithms that are sensitive to the scale of features. Here's a breakdown of its importance:

1. ****Improved Performance of Distance-Based Algorithms:****
2. ****Faster Convergence of Gradient Descent:****
3. ****Improved Regularization:****
4. ****Better Interpretation of Coefficients:****
5. ****Numerical Stability:****

```

from sklearn.base import BaseEstimator, TransformerMixin

from sklearn.pipeline import Pipeline

from sklearn.compose import ColumnTransformer

from sklearn.preprocessing import StandardScaler

# Custom transformer to add engineered attributes

class CombinedAttributesAdder(BaseEstimator, TransformerMixin):

    def __init__(self, add_bedrooms_per_room=True):

```

```

self.add_bedrooms_per_room = add_bedrooms_per_room

def fit(self, X, y=None):

    return self

def transform(self, X):

    # Assumes X is a NumPy array with the following columns:

    # total_rooms (index 3), total_bedrooms (index 2), population (index 4), households (index 5)

    rooms_per_household = X[:, 3] / X[:, 5]

    population_per_household = X[:, 4] / X[:, 5]

    if self.add_bedrooms_per_room:

        bedrooms_per_room = X[:, 2] / X[:, 3]

        return np.c_[X, rooms_per_household, population_per_household, bedrooms_per_room]

    else:

        return np.c_[X, rooms_per_household, population_per_household]

# Identify numerical and categorical columns

num_attribs = df1.drop("ocean_proximity", axis=1).columns # All numeric columns

cat_attribs = ["ocean_proximity"]

# Build numerical pipeline: impute missing values, add new attributes, then scale

num_pipeline = Pipeline([

    ('imputer', SimpleImputer(strategy="median")),

    ('attribs_adder', CombinedAttributesAdder()),

    ('std_scaler', StandardScaler()),

])

```

```
# Build the full pipeline combining numerical and categorical processing
```

```
full_pipeline = ColumnTransformer([  
    ("num", num_pipeline, num_attribs),  
    ("cat", OneHotEncoder(), cat_attribs),  
)
```

```
# Process the dataset using the pipeline
```

```
housing_prepared = full_pipeline.fit_transform(housing)  
print("Shape of processed data:", housing_prepared.shape)
```

PROGRAM 3 Implement Linear and Multi-Linear Regression algorithm using appropriate dataset

Screenshot

12/13/2025
Monday

LAB 3

x_i (week)	y_i (sales in thousands)
1	2
2	4
3	5
4	9

Ans Matrix form

$$Y = \begin{bmatrix} 2 \\ 4 \\ 5 \\ 9 \end{bmatrix}$$

$$X = \begin{bmatrix} 1 & 2 & 3 & 4 \end{bmatrix}$$

$$X = \begin{bmatrix} 1 & 1 \\ 2 & 2 \\ 3 & 3 \\ 4 & 4 \end{bmatrix}$$

$$B = \begin{bmatrix} B_0 \\ B_1 \end{bmatrix}$$

$$B = (X^T X)^{-1} X^T Y$$

$$B = \begin{bmatrix} -0.5 \\ 2.2 \end{bmatrix}$$

$$y = B_0 + B_1 x$$

$$y = -0.5 + 2.2x$$

Normal form

$$B_1 = \frac{(\sum xy) - (\bar{x})(\sum y)}{(\sum x^2) - (\bar{x})^2}$$

$$B_0 = \bar{y} - B_1 \bar{x}$$

x	y	x^2	xy
1	2	1	2
2	4	4	8
3	5	9	15
4	9	16	36

$$\bar{x} = \frac{10}{4} = 2.5$$

$$\bar{y} = \frac{20}{4} = 5$$

$$(\bar{x})^2 = 6.25$$

$$\bar{x}^2 = 7.5$$

$$\bar{xy} = 15.25$$

$$B_1 = \frac{(\sum xy) - (\bar{x})(\sum y)}{(\sum x^2) - (\bar{x})^2}$$

$$B_1 = \frac{15.25 - (2.5)(20)}{10 - 6.25}$$

$$B_1 = \frac{8.75}{3.75} = 2.33$$

$$B_0 = (\bar{y}) - B_1 \bar{x}$$

$$B_0 = 5 - 2.33(2.5)$$

$$B_0 = -0.825$$

$$y = B_0 + B_1 x$$

$$y = -0.825 + 2.33x$$

Q Diameter in Inch (in) Price in \$ (\$)

Diameter in Inch (in)	Price in \$ (\$)
8	10
10	13
12	16

Matrix method:

$$X = \begin{bmatrix} 1 & 8 \\ 1 & 10 \\ 1 & 12 \end{bmatrix}, Y = \begin{bmatrix} 10 \\ 13 \\ 16 \end{bmatrix}, B = \begin{bmatrix} B_0 \\ B_1 \end{bmatrix}$$

$$B = (X^T X)^{-1} X^T Y$$

$$B_0 = -2, B_1 = 1.5$$

$$y = B_0 + B_1 x$$

$$y = -2 + 1.5x$$

$$y = ? \text{ when } x = 20$$

$$y = -2 + 1.5(20)$$

$$y = -2 + 30$$

$$y = 28$$

Normal method

x	y	xy	x^2
8	10	80	64
10	13	130	100
12	16	192	144

$$\bar{x} = 10, \bar{y} = 13$$

$$\bar{x}^2 = 308/3$$

$$\bar{xy} = 402/3 = 134$$

$$B_1 = \frac{\sum xy - (\bar{x})(\sum y)}{(\sum x^2) - (\bar{x})^2} = \frac{134 - 130}{308/3 - 300}$$

$$B_1 = \frac{3(4)}{8} = 1.5$$

$B_1 = 1.5$

$$B_0 = (\bar{y}) - B_1 (\bar{x})$$

$$B_0 = 13 - 1.5(10)$$

$$B_0 = -2$$

$$y = B_0 + B_1 x$$

$$y = -2 + 1.5x$$

Q. Did you perform any data preprocessing steps?

Yes. Handled missing values by filling them with the column mean. Also applied label encoding to categorical (like "status") & scaled numerical features for 1000 companies.csv to normalize the data.

Q. Did you visualize the regression line for Canada-per-capita.cdv?

Yes the regression line was plotted. It shows a strong linear r/p b/w year & per capita income, meaning that as the year increases per capita also rises.

Q. To predict salary for 12 years, 10 tests com, 10 interview com

The predicted $\frac{10}{10}$

Q. Did you encode categorical variables for 1000-companies.csv?

Yes.

Code

```
# -*- coding: utf-8 -*-

import pandas as pd

import numpy as np

from sklearn import linear_model

import matplotlib.pyplot as plt


df = pd.read_csv('/content/housing_area_price.csv')

df

# Commented out IPython magic to ensure Python compatibility.

# %matplotlib inline

plt.xlabel('area')

plt.ylabel('price')

plt.scatter(df.area,df.price,color='red',marker='+')


new_df = df.drop('price',axis='columns')

new_df


price = df.price

price

# Create linear regression object

reg = linear_model.LinearRegression()

reg.fit(new_df,price)
```

```
"""(1) Predict price of a home with area = 3300 sqr ft"""
```

```
reg.predict([[3300]])
```

```
reg.coef_
```

```
reg.intercept_
```

```
"""Y = m * X + b (m is coefficient and b is intercept)"""
```

```
3300*135.78767123 + 180616.43835616432
```

```
"""(1) Predict price of a home with area = 5000 sqr ft"""
```

```
reg.predict([[5000]])
```

```
# -*- coding: utf-8 -*-
```

```
import pandas as pd
```

```
import numpy as np
```

```
from sklearn import linear_model
```

```
df = pd.read_csv('/content/homeprices_Multiple_LR.csv')
```

```
df
```

```
"""Data Preprocessing: Fill NA values with median value of a column"""
```

```
df.bedrooms.median()
```

```
df.bedrooms = df.bedrooms.fillna(df.bedrooms.median())
```

```
df
```

```
reg = linear_model.LinearRegression()
```

```
reg.fit(df.drop('price',axis='columns'),df.price)
```

```
reg.coef_
```

```
reg.intercept_
```

```
"""Find price of home with 3000 sqr ft area, 3 bedrooms, 40 year old"""
```

```
reg.predict([[3000, 3, 40]])
```

```
112.06244194*3000 + 23388.88007794*3 + -3231.71790863*40 + 221323.00186540384
```

```
import pandas as pd
```

```
from sklearn.linear_model import LinearRegression
```

```
# Load the dataset
```

```
df1 = pd.read_csv('/content/canada_per_capita_income.csv')
```



```

# Prepare the data

X = df1.year.values.reshape(-1, 1) # Features (year)

y = df1['per capita income (US$)'] # Target (per capita income)


# Create and train the linear regression model

model = LinearRegression()

model.fit(X, y)


# Predict per capita income for 2020

year_2020 = [[2020]]

predicted_income = model.predict(year_2020)


print(f'Predicted per capita income for Canada in 2020: {predicted_income[0]:.2f}')


import pandas as pd

from sklearn.linear_model import LinearRegression

import matplotlib.pyplot as plt


# Load the dataset (canada_per_capita_income.csv)

df1 = pd.read_csv('/content/canada_per_capita_income.csv')


# Prepare the data

X = df1.year.values.reshape(-1, 1) # Features (year)

```

```

y = df1['per capita income (US$)'] # Target (per capita income)

# Create and train the linear regression model

model = LinearRegression()

model.fit(X, y)

# Create the plot

plt.figure(figsize=(8, 6))

plt.scatter(X, y, color='blue', label='Data Points') # Now using the correct X and y

plt.plot(X, model.predict(X), color='red', label='Regression Line')

plt.xlabel('Year')

plt.ylabel('Per Capita Income (US$)')

plt.title('Per Capita Income in Canada over Time')

plt.legend()

plt.grid(True)

plt.show()

import pandas as pd

from sklearn.linear_model import LinearRegression

from sklearn.impute import SimpleImputer

# Load the dataset

df = pd.read_csv('/content/salary.csv')

# Prepare the data

```

```

X = df.iloc[:, :-1].values # Features (years of experience)

y = df.iloc[:, 1].values # Target (salary)


# Impute missing values with the mean

imputer = SimpleImputer(strategy='mean') # Create an imputer object with strategy as mean

X = imputer.fit_transform(X) # Fit and transform the imputer on feature data 'X'


# Create and train the linear regression model

model = LinearRegression()

model.fit(X, y)


# Predict salary for 12 years of experience

years_experience = [[12]]

predicted_salary = model.predict(years_experience)


print(f'Predicted salary for 12 years of experience: {predicted_salary[0]:.2f}')

import pandas as pd

from sklearn.linear_model import LinearRegression

from sklearn.impute import SimpleImputer


# Load the dataset

df = pd.read_csv('/content/hiring.csv')


# Handle missing values

```

```

# Convert 'experience' column to numeric, replacing non-numeric with NaN
df['experience'] = pd.to_numeric(df['experience'], errors='coerce')

imputer = SimpleImputer(strategy='mean')

df['experience'] = imputer.fit_transform(df[['experience']])
df['test_score(out of 10)'] = imputer.fit_transform(df[['test_score(out of 10)']])

# Prepare the data
X = df.drop('salary($)', axis='columns')
y = df['salary($)']

# Create and train the linear regression model
model = LinearRegression()
model.fit(X, y)

# Predict salaries for the given candidates
candidate1 = [[2, 9, 6]]
candidate2 = [[12, 10, 10]]

predicted_salary1 = model.predict(candidate1)
predicted_salary2 = model.predict(candidate2)

print(f'Predicted salary for candidate 1: ${predicted_salary1[0]:.2f}')
print(f'Predicted salary for candidate 2: ${predicted_salary2[0]:.2f}')

```

```

import pandas as pd

from sklearn.model_selection import train_test_split

from sklearn.linear_model import LinearRegression

from sklearn.preprocessing import LabelEncoder, OneHotEncoder

from sklearn.compose import ColumnTransformer


# Load the dataset

df = pd.read_csv('/content/1000_Companies.csv')


# Separate features (X) and target (y)

X = df.iloc[:, :-1].values

y = df.iloc[:, 4].values


# Encode categorical data (State)

labelencoder = LabelEncoder()

X[:, 3] = labelencoder.fit_transform(X[:, 3])


ct = ColumnTransformer(

    transformers=[('encoder', OneHotEncoder(), [3])],

    remainder='passthrough'

)

X = ct.fit_transform(X)


# Avoid dummy variable trap (remove one encoded column)

```

```
X = X[:, 1:]
```

```
# Split data into training and testing sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=0)
```

```
# Create and train the multiple linear regression model
```

```
regressor = LinearRegression()
```

```
regressor.fit(X_train, y_train)
```

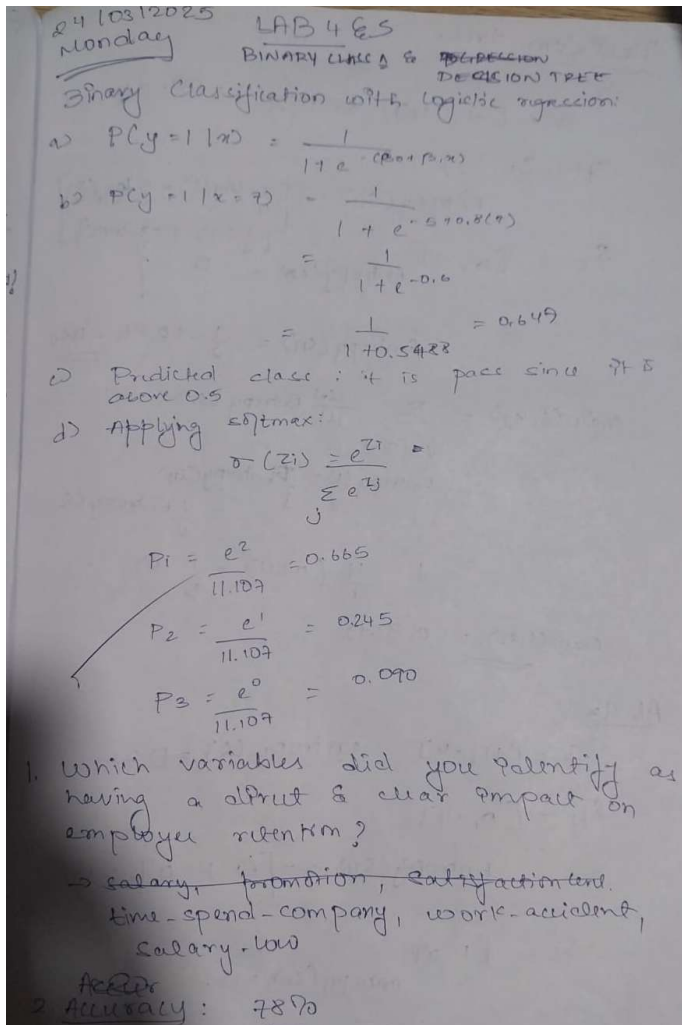
```
# Predict profit for the given values
```

```
new_prediction = regressor.predict([[1, 0, 91694.48, 515841.3, 11931.24]])
```

```
print(f'Predicted Profit: {new_prediction[0]:.2f}')
```

PROGRAM 4 Build Logistic Regression Model for a given dataset

Screenshot



Code

```
import pandas as pd

import numpy as np

df=pd.read_csv("/content/HR_comma_sep.csv")

df.head(3)

print(df.isnull().sum())

print(df.groupby('left').mean(numeric_only=True))

print(df.groupby('salary').mean(numeric_only=True))
```

```

import matplotlib.pyplot as plt

pd.crosstab(df.salary,df.left).plot(kind='bar')

plt.title('Employee Retention vs Salary')

plt.xlabel('Salary')

plt.ylabel('Number of Employees')

plt.show()

pd.crosstab(df.Department,df.left).plot(kind='bar')

plt.title('Employee Retention vs Department')

plt.xlabel('Department')

plt.ylabel('Number of Employees')

plt.show()

salary_dummies = pd.get_dummies(df.salary, prefix="salary")

dept_dummies = pd.get_dummies(df.Department, prefix="dept")


df_with_dummies = pd.concat([df, salary_dummies, dept_dummies], axis=1)


df_with_dummies = df_with_dummies.drop(['salary', 'Department'], axis=1)


X_features = ['satisfaction_level', 'last_evaluation', 'number_project', 'average_monthly_hours',
'time_spend_company', 'Work_accident', 'promotion_last_5years'] + list(salary_dummies.columns) +
list(dept_dummies.columns)

X = df_with_dummies[X_features]

y = df_with_dummies.left

```



```
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=0)


from sklearn.linear_model import LogisticRegression

model = LogisticRegression()

model.fit(X_train, y_train)


from sklearn.metrics import accuracy_score

y_pred = model.predict(X_test)


accuracy = accuracy_score(y_test, y_pred)

print("Accuracy of the model:", accuracy)
```

PROGRAM 5 Use an appropriate data set for building the decision tree (ID3) and apply this knowledge to classify a new sample.

Screenshot

Decision tree:

a₁
 $S = [+1, -4]$ Entropy(S) = 0.72
 $S_H = [+1, -3]$
 $\text{Entropy}(S_H) = -\left[\frac{1}{4} \log_2\left(\frac{1}{4}\right) + \frac{3}{4} \log_2\left(\frac{3}{4}\right)\right]$
 $= -\left[\frac{1}{4}(-2) + \frac{3}{4}(-0.415)\right]$
 $S_C = [0, -1]$ Entropy(S_C) = 0
 $\text{Entropy}(S_H) = \frac{1}{2} + 0.311 = 0.811$
 $\text{Gain}(S, a_2) = \sum_{v \in \{+1, -1\}} \frac{|S_v|}{|S|} \text{Entropy}(S_v)$
 $= \text{Entropy}(S) - \frac{4}{5} \text{Entropy}(S_H) - \frac{1}{5} \text{Entropy}(S_C)$
 $= 1 - \frac{4}{5}(0.811) - \frac{1}{5}(0)$
 $\text{Gain}(S, a_2) = 0.352$

At a₃
 $S = [+1, -4]$ entropy(S) = 0.72
 $S_H = [0, -4]$
 $\text{Entropy}(S_H) = -[0 + 0] = 0$
 $S_C = [1, 0]$ entropy(S_C) = 0
 $\text{Gain}(S, a_3) = 1$
Branching at a₃

Entropy(S) = $-\left(\frac{1}{4} \log_2\left(\frac{1}{4}\right) + \frac{3}{4} \log_2\left(\frac{3}{4}\right)\right)$
 $= 0.72$

1. 1'
 Accuracy for iris dataset $\Rightarrow 96\%$
 Confusion matrix \rightarrow

10	0	0
0	9	0
0	0	11

 The confusion matrix shows no miscalculations.

2. For petrol_consumption.csv
 regression tree predicts based on features.
 It predicts for continuous values.
 MAE : 90.60
 MCE : 16851.60
 RMSE : 129.81

24/7

Code

```
from sklearn.datasets import load_iris

from sklearn.model_selection import train_test_split

from sklearn.tree import DecisionTreeClassifier

from sklearn.metrics import accuracy_score, confusion_matrix

from sklearn import tree

import matplotlib.pyplot as plt


iris = load_iris()

X = iris.data

y = iris.target


X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)


clf = DecisionTreeClassifier()


clf.fit(X_train, y_train)


y_pred = clf.predict(X_test)


accuracy = accuracy_score(y_test, y_pred)

conf_matrix = confusion_matrix(y_test, y_pred)


print("Accuracy:", accuracy)
```

```
print("Confusion Matrix:\n", conf_matrix)
```

```
plt.figure(figsize=(12, 8))
```

```
tree.plot_tree(clf, feature_names=iris.feature_names, class_names=iris.target_names, filled=True)
```

```
plt.show()
```

```
from sklearn.datasets import load_iris
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.tree import DecisionTreeClassifier
```

```
from sklearn.metrics import accuracy_score, confusion_matrix
```

```
from sklearn import tree
```

```
import matplotlib.pyplot as plt
```

```
iris = load_iris()
```

```
X = iris.data
```

```
y = iris.target
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
clf = DecisionTreeClassifier()
```

```
clf.fit(X_train, y_train)
```

```
y_pred = clf.predict(X_test)
```

```
accuracy = accuracy_score(y_test, y_pred)
```

```
conf_matrix = confusion_matrix(y_test, y_pred)
```

```
print("Accuracy:", accuracy)
```

```
print("Confusion Matrix:\n", conf_matrix)
```

```
plt.figure(figsize=(12, 8))
```

```
tree.plot_tree(clf, feature_names=iris.feature_names, class_names=iris.target_names, filled=True)
```

```
plt.show()
```

```
import pandas as pd
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.tree import DecisionTreeRegressor
```

```
from sklearn.metrics import mean_absolute_error, mean_squared_error
```

```
import numpy as np # import numpy
```

```
data = pd.read_csv("petrol_consumption.csv")
```

```
X = data[['Petrol_tax', 'Average_income', 'Paved_Highways',
```

```
         'Population_Driver_licence(%)']]
```

```
y = data['Petrol_Consumption']
```

```

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42)

regressor = DecisionTreeRegressor()

regressor.fit(X_train, y_train)

y_pred = regressor.predict(X_test)

mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)

print("Mean Absolute Error:", mae)
print("Mean Squared Error:", mse)
print("Root Mean Squared Error:", rmse)

from sklearn.tree import plot_tree
import matplotlib.pyplot as plt

plt.figure(figsize=(15, 10))

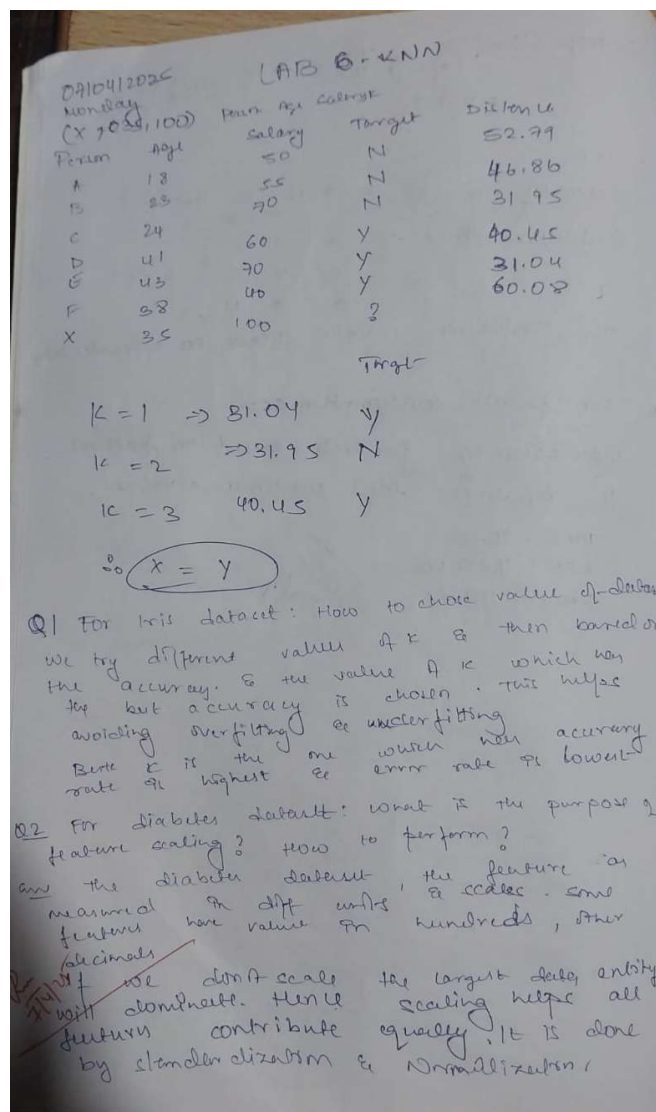
# Assuming 'data' is your original pandas DataFrame
plot_tree(regressor, feature_names=data[['Petrol_tax', 'Average_income', 'Paved_Highways',
'Population_Driver_licence(%)']].columns, filled=True, rounded=True)

plt.show()

```

PROGRAM 6 Build KNN Classification model for a given dataset.

Screenshot



Code

```
import pandas as pd

from sklearn.model_selection import train_test_split

from sklearn.neighbors import KNeighborsClassifier

from sklearn.metrics import classification_report, confusion_matrix, accuracy_score

import seaborn as sns

import matplotlib.pyplot as plt
```

```

try:

    data = pd.read_csv('/content/iris (1).csv')

except FileNotFoundError:

    print("Error: 'iris.csv' not found. Please upload the file to your Colab environment.")

    exit()


X = data.drop('species', axis=1)

y = data['species']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

knn = KNeighborsClassifier(n_neighbors=3)

knn.fit(X_train, y_train)

y_pred = knn.predict(X_test)

print("Accuracy Score:", accuracy_score(y_test, y_pred))

print("\nConfusion Matrix:")

cm = confusion_matrix(y_test, y_pred)

print(cm)


plt.figure(figsize=(8, 6))

sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',

            xticklabels=knn.classes_, yticklabels=knn.classes_)

plt.xlabel('Predicted')

plt.ylabel('Actual')

plt.title('Confusion Matrix')

```



```
plt.show()
```

```
print("\nClassification Report:")
```

```
print(classification_report(y_test, y_pred))
```

```
import pandas as pd
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.neighbors import KNeighborsClassifier
```

```
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
```

```
from sklearn.preprocessing import StandardScaler
```

```
import seaborn as sns
```

```
import matplotlib.pyplot as plt
```

```
try:
```

```
    diabetes = pd.read_csv('diabetes.csv')
```

```
except FileNotFoundError:
```

```
    print("Error: 'diabetes.csv' not found. Please ensure the file is in the current directory.")
```

```
    exit()
```

```
X = diabetes.drop('Outcome', axis=1)
```

```
y = diabetes['Outcome']
```

```
scaler = StandardScaler()
```

```
X = scaler.fit_transform(X)
```

```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

knn = KNeighborsClassifier(n_neighbors=5)

knn.fit(X_train, y_train)

y_pred = knn.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)

print(f'Accuracy: {accuracy}')
```



```

cm = confusion_matrix(y_test, y_pred)

print("Confusion Matrix:")

print(cm)

sns.heatmap(cm, annot=True, fmt="d")

plt.title('Confusion Matrix')

plt.xlabel('Predicted')

plt.ylabel('True')

plt.show()

print("Classification Report:")

print(classification_report(y_test, y_pred))
```



```

import pandas as pd

from sklearn.model_selection import train_test_split

from sklearn.neighbors import KNeighborsClassifier

from sklearn.metrics import classification_report, confusion_matrix, accuracy_score

import seaborn as sns

import matplotlib.pyplot as plt
```

```

from sklearn.preprocessing import StandardScaler

try:

    heart = pd.read_csv('heart.csv')

except FileNotFoundError:

    print("Error: 'heart.csv' not found. Please ensure the file is in the current directory.")

    exit()


X = heart.drop('target', axis=1)

y = heart['target']


scaler = StandardScaler()

X = scaler.fit_transform(X)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

best_k = 1

best_accuracy = 0


for k in range(1, 21):

    knn = KNeighborsClassifier(n_neighbors=k)

    knn.fit(X_train, y_train)

    y_pred = knn.predict(X_test)

    accuracy = accuracy_score(y_test, y_pred)

    if accuracy > best_accuracy:

        best_accuracy = accuracy

```

```

best_k = k

print(f'Best k: {best_k} with accuracy {best_accuracy}')

knn = KNeighborsClassifier(n_neighbors=best_k)

knn.fit(X_train, y_train)

y_pred = knn.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)

print(f'Accuracy: {accuracy}')
```



```

cm = confusion_matrix(y_test, y_pred)

print("Confusion Matrix:")

print(cm)
```



```

sns.heatmap(cm, annot=True, fmt="d")

plt.title('Confusion Matrix')

plt.xlabel('Predicted')

plt.ylabel('True')

plt.show()
```



```

print("Classification Report:")

print(classification_report(y_test, y_pred))
```



```

import matplotlib.pyplot as plt

import seaborn as sns
```

```
from sklearn.metrics import classification_report, confusion_matrix
```

```
cm = confusion_matrix(y_test, y_pred)
```

```
plt.figure(figsize=(8, 6))
```

```
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues")
```

```
plt.title("Confusion Matrix")
```

```
plt.xlabel("Predicted")
```

```
plt.ylabel("Actual")
```

```
plt.show()
```

```
print(classification_report(y_test, y_pred))
```

```
# prompt: For Iris dataset
```

```
# How to choose the k value? Demonstrate using accuracy rate and error
```

```
# rate. Give theory
```

```
import pandas as pd
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.neighbors import KNeighborsClassifier
```

```
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
```

```
import seaborn as sns
```

```
import matplotlib.pyplot as plt
```

```
from sklearn.preprocessing import StandardScaler
```

```
# Load the Iris dataset
```

try:

```
data = pd.read_csv('/content/iris (1).csv')
```

except FileNotFoundError:

```
print("Error: 'iris (1).csv' not found. Please upload the file to your Colab environment.")
```

```
exit()
```

Prepare the data

```
X = data.drop('species', axis=1)
```

```
y = data['species']
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Scale the data (important for KNN)

```
scaler = StandardScaler()
```

```
X_train = scaler.fit_transform(X_train)
```

```
X_test = scaler.transform(X_test)
```

Find the optimal k value

```
error_rates = []
```

```
for k in range(1, 31): # Test k values from 1 to 30
```

```
    knn = KNeighborsClassifier(n_neighbors=k)
```

```
    knn.fit(X_train, y_train)
```

```
    y_pred = knn.predict(X_test)
```

```
    error_rates.append(1 - accuracy_score(y_test, y_pred)) # Error rate = 1 - accuracy
```

```

# Plot error rates

plt.figure(figsize=(10, 6))

plt.plot(range(1, 31), error_rates, color='blue', linestyle='dashed', marker='o',
         markerfacecolor='red', markersize=10)

plt.title('Error Rate vs. K Value')

plt.xlabel('K')

plt.ylabel('Error Rate')

plt.show()


# Theory for choosing k:

# The optimal 'k' value minimizes the error rate.

# Very small k (e.g., 1) can lead to overfitting, being too sensitive to noise.

# Very large k (e.g., 30) can lead to underfitting, smoothing out the decision boundaries too much.

# We seek a k that balances these extremes, as shown by the error rate plot.


#Select k based on the minimum error rate observed in the plot

best_k = error_rates.index(min(error_rates)) + 1 #Add 1 as the index starts from 0


# Train and evaluate the model with the best k

knn = KNeighborsClassifier(n_neighbors=best_k)

knn.fit(X_train, y_train)

y_pred = knn.predict(X_test)

```

```

# Evaluate the model

print("Accuracy Score:", accuracy_score(y_test, y_pred))

print("\nConfusion Matrix:")

cm = confusion_matrix(y_test, y_pred)

print(cm)

print("\nClassification Report:")

print(classification_report(y_test, y_pred))


plt.figure(figsize=(8, 6))

sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',

            xticklabels=knn.classes_, yticklabels=knn.classes_)

plt.xlabel('Predicted')

plt.ylabel('Actual')

plt.title('Confusion Matrix')

plt.show()


import pandas as pd

from sklearn.model_selection import train_test_split

from sklearn.neighbors import KNeighborsClassifier

from sklearn.metrics import accuracy_score

import matplotlib.pyplot as plt


# Load data

```



```

df = pd.read_csv('/content/iris (1).csv')

X = df.iloc[:, :-1]

y = df.iloc[:, -1]

# Train-test split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=1)

# Store accuracy and error rate

accuracy = []

error_rate = []

# Try k from 1 to 20

for k in range(1, 21):

    knn = KNeighborsClassifier(n_neighbors=k)

    knn.fit(X_train, y_train)

    preds = knn.predict(X_test)

    acc = accuracy_score(y_test, preds)

    accuracy.append(acc)

    error_rate.append(1 - acc)

# Plot

plt.figure(figsize=(10,5))

plt.plot(range(1, 21), accuracy, label='Accuracy')

plt.plot(range(1, 21), error_rate, label='Error Rate')

```

```
plt.xlabel('K Value')

plt.ylabel('Rate')

plt.title('K vs Accuracy and Error Rate')

plt.legend()

plt.show()


import pandas as pd

from sklearn.preprocessing import StandardScaler


# Load data

df = pd.read_csv('/content/diabetes.csv')

X = df.drop('Outcome', axis=1) # Features

y = df['Outcome']           # Target


# Perform scaling

scaler = StandardScaler()

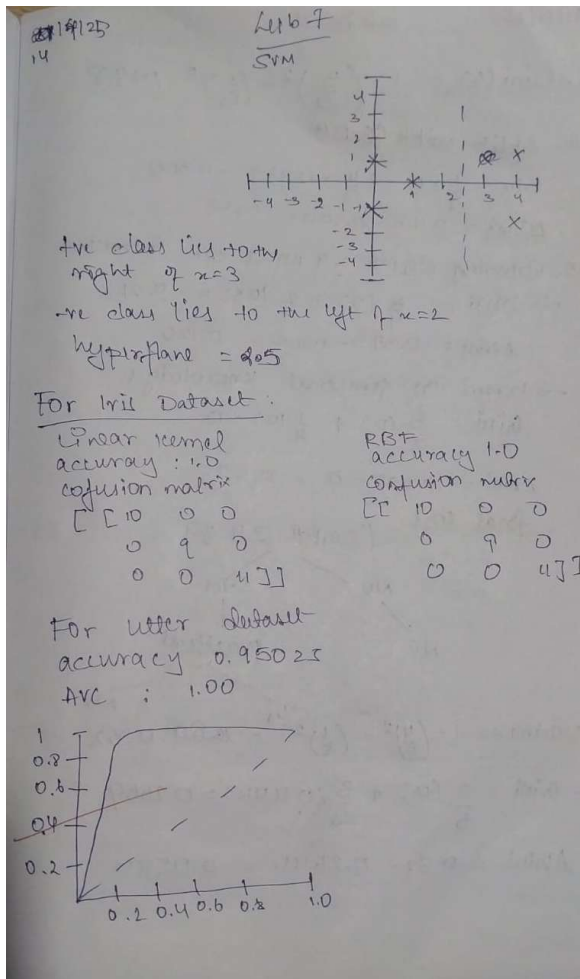
X_scaled = scaler.fit_transform(X)


# Convert back to DataFrame (optional)

X_scaled_df = pd.DataFrame(X_scaled, columns=X.columns)
```

PROGRAM 7 Build Support vector machine model for a given dataset

Screenshot



Code

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
positive_class = np.array([[4, 1], [4, -1], [6, 0]])
```

```
negative_class = np.array([[1, 0], [0, 1], [0, -1]])
```

```
plt.figure(figsize=(8, 6))
```

```

plt.scatter(positive_class[:, 0], positive_class[:, 1], color='red', label='Positive Class', s=100,
edgecolors='black')

plt.scatter(negative_class[:, 0], negative_class[:, 1], color='blue', label='Negative Class', s=100,
edgecolors='black')

all_points = np.concatenate([positive_class, negative_class])

labels = ["(4,1)", "(4,-1)", "(6,0)", "(1,0)", "(0,1)", "(0,-1)"]

for i, txt in enumerate(labels):

    plt.annotate(txt, (all_points[i][0], all_points[i][1]), textcoords="offset points", xytext=(0,5),
ha='center', fontsize=10)

x_values = np.linspace(-1, 7, 100)

y_values = np.zeros_like(x_values)

plt.plot(x_values, y_values, color='black', linestyle='--', label='Optimal Hyperplane (y = 0)')

plt.plot(x_values, y_values + 1, color='gray', linestyle=':', label='Margin at y = 1')

plt.plot(x_values, y_values - 1, color='gray', linestyle=':', label='Margin at y = -1')

plt.title('Optimal Hyperplane for SVM (Visual Approximation)', fontsize=14)

plt.xlabel('x1')

plt.ylabel('x2')

plt.xlim(-1, 7)

plt.ylim(-2, 2)

plt.axhline(0, color='black', linewidth=0.5)

```

```

plt.axvline(0, color='black',linewidth=0.5)

plt.legend()


plt.grid(True)

plt.show()

import pandas as pd

from sklearn.model_selection import train_test_split

from sklearn.svm import SVC

from sklearn.metrics import accuracy_score, confusion_matrix

import seaborn as sns

import matplotlib.pyplot as plt


data = pd.read_csv('/content/iris (1) (1).csv')


X = data.drop('species', axis=1)

y = data['species']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)


svm_rbf = SVC(kernel='rbf')

svm_rbf.fit(X_train, y_train)

y_pred_rbf = svm_rbf.predict(X_test)

accuracy_rbf = accuracy_score(y_test, y_pred_rbf)

cm_rbf = confusion_matrix(y_test, y_pred_rbf)

```

```

print("SVM with RBF Kernel:")

print("Accuracy:", accuracy_rbf)

print("Confusion Matrix:\n", cm_rbf)


plt.figure(figsize=(6, 4))

sns.heatmap(cm_rbf, annot=True, fmt='d', cmap='Blues',

            xticklabels=data['species'].unique(),

            yticklabels=data['species'].unique())

plt.xlabel('Predicted')

plt.ylabel('Actual')

plt.title('Confusion Matrix (RBF Kernel)')

plt.show()

```

```

svm_linear = SVC(kernel='linear')

svm_linear.fit(X_train, y_train)

y_pred_linear = svm_linear.predict(X_test)

accuracy_linear = accuracy_score(y_test, y_pred_linear)

cm_linear = confusion_matrix(y_test, y_pred_linear)

```

```

print("\nSVM with Linear Kernel:")

print("Accuracy:", accuracy_linear)

print("Confusion Matrix:\n", cm_linear)

```

```

plt.figure(figsize=(6, 4))

```

```

sns.heatmap(cm_linear, annot=True, fmt='d', cmap='Blues',
            xticklabels=data['species'].unique(),
            yticklabels=data['species'].unique())

plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix (Linear Kernel)')
plt.show()

import numpy as np

import matplotlib.pyplot as plt

import pandas as pd

from sklearn.model_selection import train_test_split

from sklearn.svm import SVC

from sklearn.metrics import accuracy_score, confusion_matrix, roc_curve, auc

import seaborn as sns

from sklearn.preprocessing import label_binarize

from sklearn.multiclass import OneVsRestClassifier

data = pd.read_csv('/content/letter-recognition.csv') # Replace with the correct path if necessary

X = data.drop('letter', axis=1)

y = data['letter']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

```

```
svm_classifier = SVC(kernel='rbf', probability=True) # probability=True is needed for ROC curve
```

```
svm_classifier.fit(X_train, y_train)
```

```
y_pred = svm_classifier.predict(X_test)
```

```
accuracy = accuracy_score(y_test, y_pred)
```

```
cm = confusion_matrix(y_test, y_pred)
```

```
print("SVM Classifier:")
```

```
print("Accuracy:", accuracy)
```

```
print("Confusion Matrix:\n", cm)
```

```
plt.figure(figsize=(10, 8))
```

```
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=np.unique(y),  
yticklabels=np.unique(y))
```

```
plt.xlabel('Predicted')
```

```
plt.ylabel('Actual')
```

```
plt.title('Confusion Matrix')
```

```
plt.show()
```

```
y_test_bin = label_binarize(y_test, classes=np.unique(y))
```

```
n_classes = y_test_bin.shape[1]
```

```
classifier = OneVsRestClassifier(SVC(kernel='rbf', probability=True))
```

```
classifier.fit(X_train, y_train)
```



```

y_score = classifier.predict_proba(X_test)

fpr = dict()
tpr = dict()
roc_auc = dict()
for i in range(n_classes):
    fpr[i], tpr[i], _ = roc_curve(y_test_bin[:, i], y_score[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])

fpr["micro"], tpr["micro"], _ = roc_curve(y_test_bin.ravel(), y_score.ravel())
roc_auc["micro"] = auc(fpr["micro"], tpr["micro"])

plt.figure(figsize=(8, 6))
plt.plot(fpr["micro"], tpr["micro"],
         label='micro-average ROC curve (area = {0:0.2f})'
         .format(roc_auc["micro"]))
plt.plot([0, 1], [0, 1], 'k--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Micro-averaged ROC Curve')
plt.legend(loc="lower right")
plt.show()

print(f'Micro-averaged AUC: {roc_auc["micro"]}')

```

PROGRAM 8 Implement Random forest ensemble method on a given dataset.

Screenshot

6/5/25 Lab 8

1. Gini (A) = $1 - \left(\frac{3}{5}\right)^2 - \left(\frac{2}{5}\right)^2 = 0.48$

2. Split with CnPA

$$Gini = \frac{1}{5} \times 0 + \frac{4}{5} \times 0.375 = 0.300$$

$$\Delta Gini = 0.48 - 0.300 = 0.18$$

3. Growing CnPA ≥ 9 nodes on Interaction

$$\rightarrow Gini = \frac{2}{9} (0) + \frac{7}{9} (0.5) = 0.21$$

$$\Delta Gini = 0.375 - 0.21 = 0.165$$

\rightarrow based on practical knowledge

$$Gini = \frac{3}{9} (0) + \frac{1}{4} (0.5) = 0$$

$$\Delta Gini = 0.375 - 0 = 0.375$$

Final Cost [CnPA ≥ 9]

```

graph TD
    A[ ] -- NO --> B[ ]
    A -- YES --> C[Practical]
    B -- NO --> D[ ]
    C -- YES --> E[ ]
    C -- NO --> F[ ]
    
```

1. Gini = $1 - \left(\frac{4}{5}\right)^2 - \left(\frac{1}{5}\right)^2 = 0.32$

2. Gini = $\frac{2}{5} (0) + \frac{3}{5} (0.444) = 0.2664$

$$\Delta Gini = 0.32 - 0.2664 = 0.0536$$

3. Growing the "no" ~~base~~ with other features.

\rightarrow CnPA

$$\text{weighted Gini} : \frac{1}{3} (0.5) + \frac{2}{3} (0.5) = 0.333$$

$$\Delta Gini = 0.111$$

\rightarrow Practical knowledge:

$$\text{weighted Gini} = 0$$

$$\Delta Gini = 0.444$$

```

graph TD
    A[Interactive] -- YES --> B[ ]
    A -- NO --> C[Practical knowledge]
    B -- YES --> D[ ]
    C -- Good? --> E[ ]
    C -- Avg --> F[ ]
    E -- YES --> G[ ]
    F -- NO --> H[ ]
    
```

RF	n Estimators	mean accuracy
1	10	0.9667
2	50	0.9667
3	100	0.9667
4	150	0.9667
5	200	0.9667

The best can be chosen with n estimators

Code

```
import pandas as pd

from sklearn.model_selection import train_test_split

from sklearn.ensemble import RandomForestClassifier

from sklearn.metrics import accuracy_score

import matplotlib.pyplot as plt


# Load the dataset

df = pd.read_csv('/content/iris (1).csv')


# Prepare features and target

X = df.drop(columns=['species']) # Assuming 'species' is the target column
y = df['species']


# Split into training and testing sets

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)


# Build Random Forest with default n_estimators (10)

rf_default = RandomForestClassifier(n_estimators=10, random_state=42)

rf_default.fit(X_train, y_train)

y_pred_default = rf_default.predict(X_test)


# Measure accuracy

default_score = accuracy_score(y_test, y_pred_default)
```

```

print(f'Default RF accuracy (n_estimators=10): {default_score:.4f}')

# Fine-tune the number of trees

scores = []

n_range = range(1, 101)

for n in n_range:

    rf = RandomForestClassifier(n_estimators=n, random_state=42)

    rf.fit(X_train, y_train)

    y_pred = rf.predict(X_test)

    score = accuracy_score(y_test, y_pred)

    scores.append(score)

# Find the best score and number of trees

best_score = max(scores)

best_n = n_range[scores.index(best_score)]

print(f'Best RF accuracy: {best_score:.4f} with n_estimators={best_n}')

# Optional: Plot accuracy vs number of estimators

plt.figure(figsize=(10, 6))

plt.plot(n_range, scores, marker='o')

plt.title('Random Forest Accuracy vs Number of Trees')

plt.xlabel('Number of Trees (n_estimators)')

plt.ylabel('Accuracy')

plt.grid(True)

plt.show()

```

PROGRAM 9 Implement Boosting ensemble method on a given dataset.

Screenshot

5/5/20

L	GAPP	Actual Job	y_b
29	+1	y_b	
29	-1	y_b	
29	-1	y_b	
29	+1	y_b	
29	+1	y_b	

$E_{GAPP} = 4 \times \frac{1}{6} = 0.66$
 $\alpha_{GAPP} = 0.347 = \frac{1}{2} \ln \left(\frac{1-0.333}{0.333} \right)$
 $Z_{GAPP} = \frac{1}{6} \times 4 \times e^{-0.347} + \frac{1}{6} \times 2 \times e^{0.347}$
 $= 0.9428$
 $w_{t+1} = \frac{\frac{1}{6} \times e^{-0.347}}{0.9428} = 0.1249$
 $w_{t-1} = \frac{\frac{1}{6} \times e^{0.347}}{0.9428} = 0.2501$

Interactions

Actual Job	y_b	Weight
Yes	y_b	0.1249
No	y_b	0.2501
No	No	0.2501
No	No	0.1249
Yes	y_b	0.1249
Yes	y_b	0.1249

5/5/20

$E_{interactions} \sim 0.2501$
 $\alpha = \frac{1}{2} \ln \left[\frac{1-0.2501}{0.2501} \right] = 0.2190$
 $E_{interactions} = 0.1249 \times 4 \times e^{-0.199} + 0.2501 \times 1 \times e^{-0.949} + 0.2501 \times 1 \times e^{0.241}$
 $= 0.866$
 $w_{t+1} = \frac{0.2501 \times e^{-0.549}}{0.866} = 0.1664$
 $w_{t-1} = \frac{0.2501 \times e^{-0.549}}{0.866} = 0.5001$

No pattern or knowledge

P_t misclassified by practical

Communication Skill	Actual Job	Weight
Good	Yes	0.0532
Moderate	Yes	0.5001
Moderate	No	0.1662
Good	No	0.0831
Moderate	Yes	0.0532
Moderate	Yes	0.0832

P
5/5/20

Code

```
import pandas as pd

import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split

from sklearn.preprocessing import LabelEncoder

from sklearn.ensemble import AdaBoostClassifier

from sklearn.metrics import accuracy_score

from sklearn.tree import DecisionTreeClassifier
```

```

# Load dataset

df = pd.read_csv("/content/income.csv")

# Drop rows with missing values

df.dropna(inplace=True)

# Encode categorical columns

label_encoders = {}

for column in df.select_dtypes(include=['object']).columns:

    le = LabelEncoder()

    df[column] = le.fit_transform(df[column])

    label_encoders[column] = le

# Separate features and target

X = df.drop(columns=['income_level'], errors='ignore', axis=1)

y = df['income_level']

# Split into training and testing sets

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# AdaBoost with 10 estimators

model_10 = AdaBoostClassifier(n_estimators=10, random_state=42)

model_10.fit(X_train, y_train)

y_pred_10 = model_10.predict(X_test)

score_10 = accuracy_score(y_test, y_pred_10)

print(f'Accuracy with 10 estimators: {score_10:.4f}')

# Fine-tune number of estimators

best_score = 0

```

```

best_n = 0

estimators_range = list(range(10, 201, 10))

scores = []

for n in estimators_range:

    model = AdaBoostClassifier(n_estimators=n, random_state=42)

    model.fit(X_train, y_train)

    y_pred = model.predict(X_test)

    score = accuracy_score(y_test, y_pred)

    scores.append(score)

    print(f'n_estimators={n}, Accuracy={score:.4f}')

    if score > best_score:

        best_score = score

        best_n = n

print(f'\nBest Accuracy: {best_score:.4f} using {best_n} estimators')

# Plot accuracy vs number of estimators

plt.figure(figsize=(7, 4))

plt.plot(estimators_range, scores, marker='o', linestyle='-', color='blue')

plt.title("Accuracy vs Number of Estimators (AdaBoost)")

plt.xlabel("Number of Estimators (Trees)")

plt.ylabel("Accuracy")

plt.grid(True)

plt.xticks(estimators_range)

plt.tight_layout()

plt.show()

```

PROGRAM 10 Build k-Means algorithm to cluster a set of data stored in a .CSV file.

Screenshot

12/05/25 LAB 10 means

Q. Compute two clusters using k-means algorithm for clustering whose cluster centers are $(1, 1)$ & $(5, 7)$. Execute for 2 iterations.

→ Iteration 1

Record	Euclidean dist to $C_1(1,1)$	to $C_2(5,7)$	Assignment
$(1, 1)$	0	7.41	C_1
$(1.5, 2)$	1.12	6.12	C_1
$(3, 4)$	2.61	3.61	C_1
$(5, 7)$	7.21	0.0	C_2
$(3.5, 5)$	4.12	2.5	C_2
$(4.5, 5)$	5.31	2.06	C_2
$(3.5, 4.5)$	4.20	2.92	C_2

New centroids:

$$C_1 = \frac{1+1.5+3}{3}, \frac{1+2+4}{3} = 1.83, 2.33$$

$$C_2 = \frac{5+3.5+4.5+3.5}{4}, \frac{7+5+5+4.5}{4} = 4.12, 5.37$$

Iteration 2

Record	Dist to $C_1(1.83, 2.33)$	Dist to $C_2(4.12, 5.37)$	Assignment
$(1, 1)$	1.67	5.37	C_1
$(1.5, 2)$	0.47	4.27	C_1
$(3, 4)$	2.04	1.77	C_2
$(5, 7)$	5.64	1.85	C_2
$(3.5, 5)$	3.15	0.72	C_2
$(4.5, 5)$	3.73	0.53	C_2
$(3.5, 4.5)$	2.74	1.07	C_2

∴ New Clusters are
 $C_1 = \{A_1, B_2\}$, $C_2 = \{B_3, B_4, B_5, B_6\}$

New Centroids
 $C_1 = \frac{2.5}{2}, \frac{3}{2}$ | $C_2 = \frac{19.5}{5}, \frac{25.5}{5}$

Q. For Iris dataset:

The elbow plot (inertia vs k) shows a sharp elbow at $k=3$, indicating that three clusters is the optimal choice for the dataset / width of the iris.

Code

```
import pandas as pd

import numpy as np

from sklearn.model_selection import train_test_split

from sklearn.preprocessing import StandardScaler

from sklearn.cluster import KMeans

from sklearn.metrics import accuracy_score
```



```
from scipy.stats import mode

import matplotlib.pyplot as plt


# Step 1: Generate sample data and save to CSV

np.random.seed(42)

names = [f"Person_{i}" for i in range(50)]

ages = np.random.randint(20, 60, 50)

income = np.random.randint(30000, 120000, 50)


df = pd.DataFrame({'Name': names, 'Age': ages, 'Income': income})

df.to_csv("income.csv", index=False)


# Step 2: Load the data

data = pd.read_csv("income.csv")


# Drop 'Name' and extract features

X = data[['Age', 'Income']]


# Step 3: Split the data

X_train, X_test = train_test_split(X, test_size=0.2, random_state=42)


# Step 4: Perform scaling

scaler = StandardScaler()

X_train_scaled = scaler.fit_transform(X_train)
```

```

X_test_scaled = scaler.transform(X_test)

# Step 5: Plot SSE vs number of clusters (Elbow method)

sse = []

k_range = range(1, 11)

for k in k_range:

    kmeans = KMeans(n_clusters=k, random_state=42)

    kmeans.fit(X_train_scaled)

    sse.append(kmeans.inertia_)


plt.figure(figsize=(8, 4))

plt.plot(k_range, sse, marker='o')

plt.xlabel('Number of clusters')

plt.ylabel('SSE (Inertia)')

plt.title('Elbow Method For Optimal k')

plt.grid(True)

plt.show()


# Step 6: Choose optimal number of clusters (say 3) and fit model

optimal_k = 3

kmeans = KMeans(n_clusters=optimal_k, random_state=42)

kmeans.fit(X_train_scaled)


# Predict on test data

```

```

predictions = kmeans.predict(X_test_scaled)

# Note: There's no ground truth labels, but for demonstration,
# we can try assigning true clusters (via KMeans on full data)
# and see if predicted clusters align

# Fit on full data to assign pseudo-labels
full_kmeans = KMeans(n_clusters=optimal_k, random_state=42)
true_clusters = full_kmeans.fit_predict(scaler.fit_transform(X))

# Align predicted clusters using majority voting (only for demonstration)
# Match predicted labels to closest true labels
def map_clusters(true_labels, pred_labels):
    labels = np.zeros_like(pred_labels)

    for i in range(optimal_k):
        mask = (pred_labels == i)
        if np.sum(mask) == 0:
            continue

        labels[mask] = mode(true_labels[mask])[0]

    return labels

mapped_preds = map_clusters(true_clusters[X_test.index], predictions)
accuracy = accuracy_score(true_clusters[X_test.index], mapped_preds)
print(f'Approximate Clustering Accuracy: {accuracy:.2f}')

```

```

import numpy as np

import pandas as pd

import matplotlib.pyplot as plt

from sklearn.datasets import load_iris

from sklearn.cluster import KMeans

from sklearn.preprocessing import StandardScaler

from sklearn.metrics import silhouette_score


# Step 1: Load Iris dataset

iris = load_iris()

df = pd.DataFrame(iris.data, columns=iris.feature_names)

df['target'] = iris.target


# Keep only petal length and petal width

X = df[['petal length (cm)', 'petal width (cm)']].values


# Step 2: Check impact of scaling

# Try without scaling

sse_unscaled = []

for k in range(1, 11):

    kmeans = KMeans(n_clusters=k, random_state=42)

    kmeans.fit(X)

    sse_unscaled.append(kmeans.inertia_)

```

```

# Now scale the features

scaler = StandardScaler()

X_scaled = scaler.fit_transform(X)


sse_scaled = []

for k in range(1, 11):

    kmeans = KMeans(n_clusters=k, random_state=42)

    kmeans.fit(X_scaled)

    sse_scaled.append(kmeans.inertia_)


# Step 3: Plot Elbow Comparison (Scaled vs Unscaled)

plt.figure(figsize=(10, 5))


plt.plot(range(1, 11), sse_unscaled, marker='o', label='Unscaled')
plt.plot(range(1, 11), sse_scaled, marker='s', label='Scaled')

plt.title('Elbow Method (Petal Features Only)')

plt.xlabel('Number of Clusters (k)')

plt.ylabel('SSE (Inertia)')

plt.legend()

plt.grid(True)

plt.show()

```

PROGRAM 11 Implement Dimensionality reduction using Principal Component Analysis (PCA) method.

Screenshot

R100120 Lab 11
 PCA
 8. Perform dimensionality reduction from 4 to 1

Standardize the dataset
 $\bar{x}_1 = \frac{4+8+13+22}{4} = 12$
 $\bar{x}_2 = \frac{11+4+5+14}{4} = 8.5$
 $X_1 = \begin{bmatrix} 4-12 & 8-12 & 13-12 & 22-12 \\ 11-8.5 & 4-8.5 & 5-8.5 & 14-8.5 \end{bmatrix}$
 $X_1 = \begin{bmatrix} -8 & -4 & 1 & 10 \\ 2.5 & -4.5 & -3.5 & 5.5 \end{bmatrix}$
 $X_1^T = \begin{bmatrix} -8 & 0 & 5 & -1 \\ 2.5 & -4.5 & -3.5 & 5.5 \end{bmatrix}$

Covariance matrix
 $C = \frac{1}{n-1} X_1^T X_1$
 $C = \frac{1}{3} \begin{bmatrix} -8 & 0 & 5 & -1 \\ 2.5 & -4.5 & -3.5 & 5.5 \end{bmatrix} \begin{bmatrix} -8 & 2.5 \\ 0 & -4.5 \\ 5 & -3.5 \\ -1 & 5.5 \end{bmatrix}$
 $C = \frac{1}{3} \begin{bmatrix} 42 & -33 \\ -33 & 69 \end{bmatrix}$
 $C = \begin{bmatrix} 14 & -11 \\ -11 & 23 \end{bmatrix}$
 $\det(C - \lambda I) = 0$
 $\begin{vmatrix} 14-\lambda & -11 \\ -11 & 23-\lambda \end{vmatrix} = 0$
 $\lambda^2 - 37\lambda + 201 = 0$

$\lambda = 30.3849$
 Consider eigenvector $\begin{bmatrix} x \\ y \end{bmatrix}$
 $\begin{bmatrix} 14-30.3849 & -11 \\ -11 & 23-30.3849 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$
 $\begin{bmatrix} -16.3849x & -11y \\ -11x & -7.3849y \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$
 $-16.3849x - 11y = 0$
 $x = -0.6713y$
 taking $y = 1$
 $\begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} -0.6713 \\ 1 \end{bmatrix}$
 Normalized
 $x = \frac{-0.6713}{\sqrt{(-0.6713)^2 + 1^2}}, y = \frac{1}{\sqrt{(-0.6713)^2 + 1^2}}$
 $C_1 = \begin{bmatrix} -0.5574 \\ 0.8303 \end{bmatrix}$
 Calculate principal component:
 $Z = X_1 C_1 = \begin{bmatrix} -4 & 2.5 \\ 0 & -4.5 \\ 5 & -3.5 \\ -1 & 5.5 \end{bmatrix} \begin{bmatrix} -0.5574 \\ 0.8303 \end{bmatrix}$
 $Z = \begin{bmatrix} +4.305 \\ -3.736 \\ -5.693 \\ +5.124 \end{bmatrix}$

For Heart.csv dataset
 Accuracy before PCA
 Logistic regression: 0.9016
 SVM: 0.8525
 Random forest: 0.8361
 Accuracy After PCA
 Logistic: 0.8659
 SVM: 0.8689
 Random forest: 0.8361

Code

```

import pandas as pd

from sklearn.preprocessing import LabelEncoder, OneHotEncoder, StandardScaler

from sklearn.compose import ColumnTransformer

from sklearn.pipeline import Pipeline

from sklearn.model_selection import train_test_split

from sklearn.svm import SVC
  
```

```

from sklearn.linear_model import LogisticRegression

from sklearn.ensemble import RandomForestClassifier

from sklearn.decomposition import PCA

from sklearn.metrics import accuracy_score


# 1. Load data

df = pd.read_csv("heart.csv")


# 2. Label-encode binary text columns

le = LabelEncoder()

for col in ["Sex", "ExerciseAngina"]:

    df[col] = le.fit_transform(df[col])


# 3. Separate features and target

X = df.drop("HeartDisease", axis=1)

y = df["HeartDisease"]


# 4. Build preprocessing pipeline:

# - One-hot for multi-category columns (using sparse_output=False)

# - passthrough the rest

# - then scale everything

cat_cols = ["ChestPainType", "RestingECG", "ST_Slope"]

preprocessor = Pipeline([

    ("onehot", ColumnTransformer([

```

```

        ("ohe", OneHotEncoder(sparse_output=False, drop="first"), cat_cols)
    ], remainder="passthrough")),
    ("scaler", StandardScaler())
])

```

5. Apply preprocessing

```
X_proc = preprocessor.fit_transform(X)
```

6. Train/test split

```

X_train, X_test, y_train, y_test = train_test_split(
    X_proc, y, test_size=0.2, random_state=42
)

```

7. Define models

```

models = {
    "SVM": SVC(random_state=42),
    "LogisticRegression": LogisticRegression(max_iter=1000, random_state=42),
    "RandomForest": RandomForestClassifier(random_state=42)
}

```

8. Train & evaluate before PCA

```

print("=== Accuracies BEFORE PCA ===")

scores_before = {}

for name, clf in models.items():

```



```

clf.fit(X_train, y_train)

preds = clf.predict(X_test)

acc = accuracy_score(y_test, preds)

scores_before[name] = acc

print(f'{name:17s}: {acc:.4f}')

```

9. Apply PCA (retain 95% variance)

```

pca = PCA(n_components=0.95, random_state=42)

X_train_pca = pca.fit_transform(X_train)

X_test_pca = pca.transform(X_test)

print(f'\nPCA retained {pca.n_components_} components, "

    f'explained variance = {pca.explained_variance_ratio_.sum():.4f}\n')

```

10. Train & evaluate after PCA

```

print("=== Accuracies AFTER PCA ===")

scores_after = {}

for name, clf in models.items():

    clf.fit(X_train_pca, y_train)

    preds = clf.predict(X_test_pca)

    acc = accuracy_score(y_test, preds)

    scores_after[name] = acc

print(f'{name:17s}: {acc:.4f}')

```